

CS450 Computer Networks

The slides used in class are derived from the slides available on
our text book companion website:

http://wps.pearsoned.com/ecs_kurose_compnets_6/
copyright 1996-2012 J.F Kurose and K.W. Ross

© 2012 Maharishi University of Management

All additional course materials are copyright protected by international copyright laws and remain the property of the Maharishi University of Management. The materials are accessible only for the personal use of students enrolled in this course and only for the duration of the course. Any copying and distributing are not allowed and subject to legal action.

CS450 Computer Networks

Lesson 9

Transport Layer - Reliable Data Transfer

Purification Leads to Progress

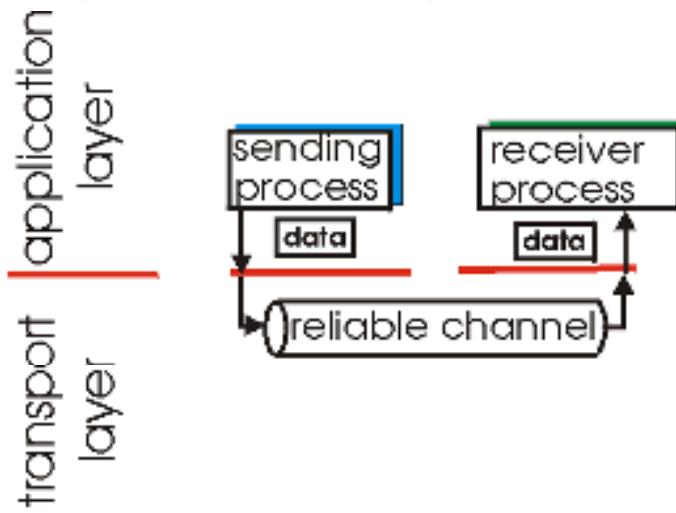
Lesson 9: Transport Layer - Reliable Data Transfer

Our goal:

- ❖ understand principles behind reliable data transfer as a transport layer service

Principles of Reliable data transfer

- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!

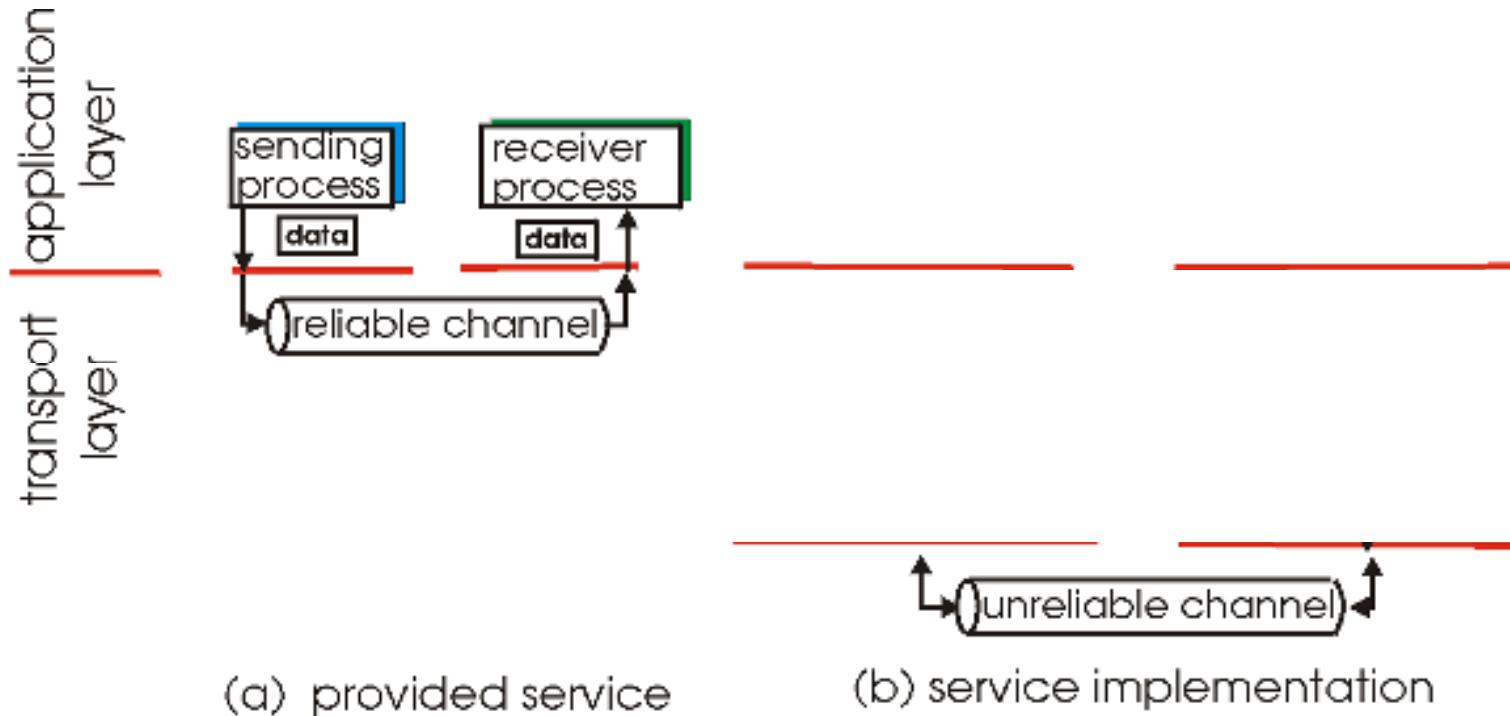


(a) provided service

- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of Reliable data transfer

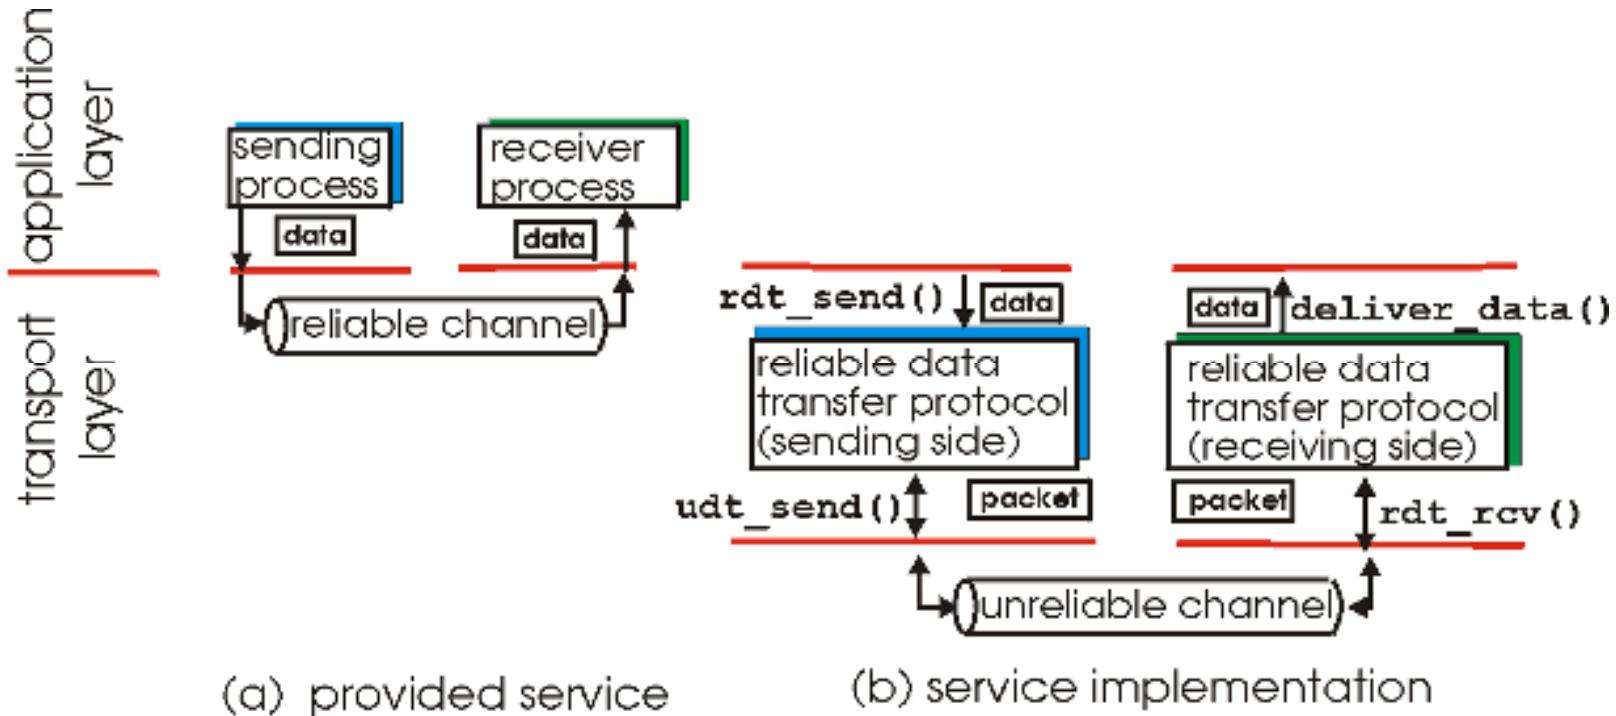
- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of Reliable data transfer

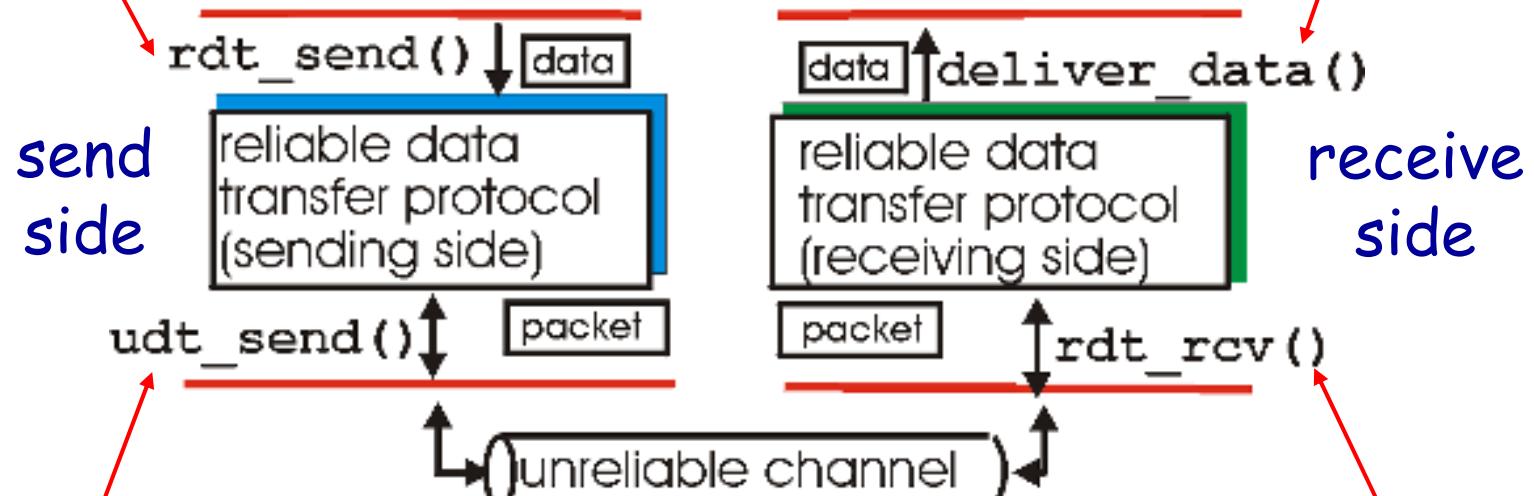
- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: getting started

rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer



udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver

deliver_data() : called by
rdt to deliver data to upper

rdt_rcv() : called when packet
arrives on rcv-side of channel

Reliable data transfer: getting started

We will:

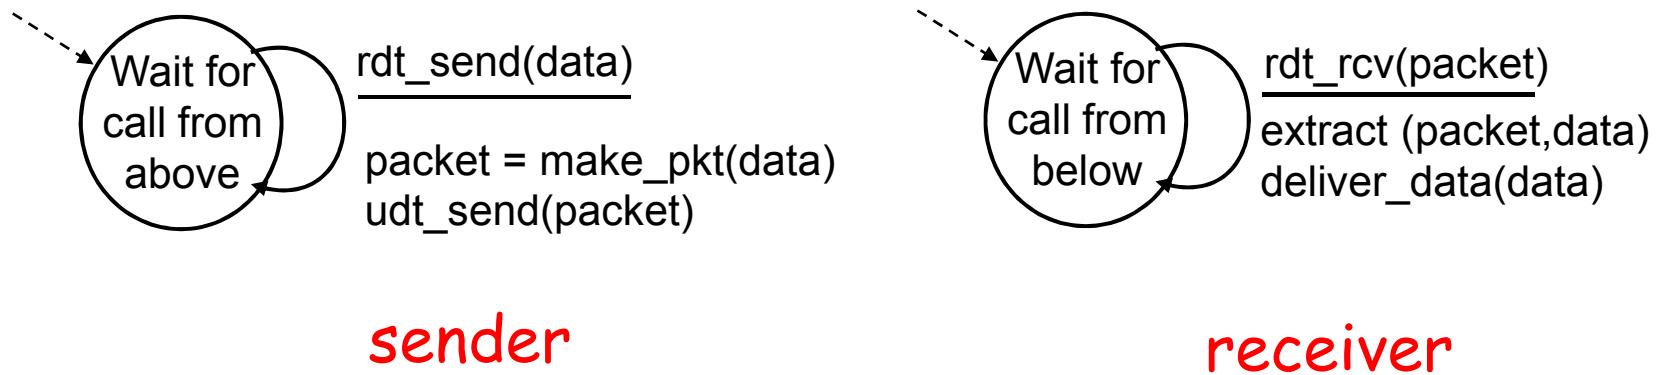
- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
 - but control info will flow on both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver

state: when in this "state" next state uniquely determined by next event



Rdt1.0: reliable transfer over a reliable channel

- ❖ underlying channel perfectly reliable
 - no bit errors
 - no loss of packets
- ❖ separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver read data from underlying channel



Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ *the question:* how to recover from errors:

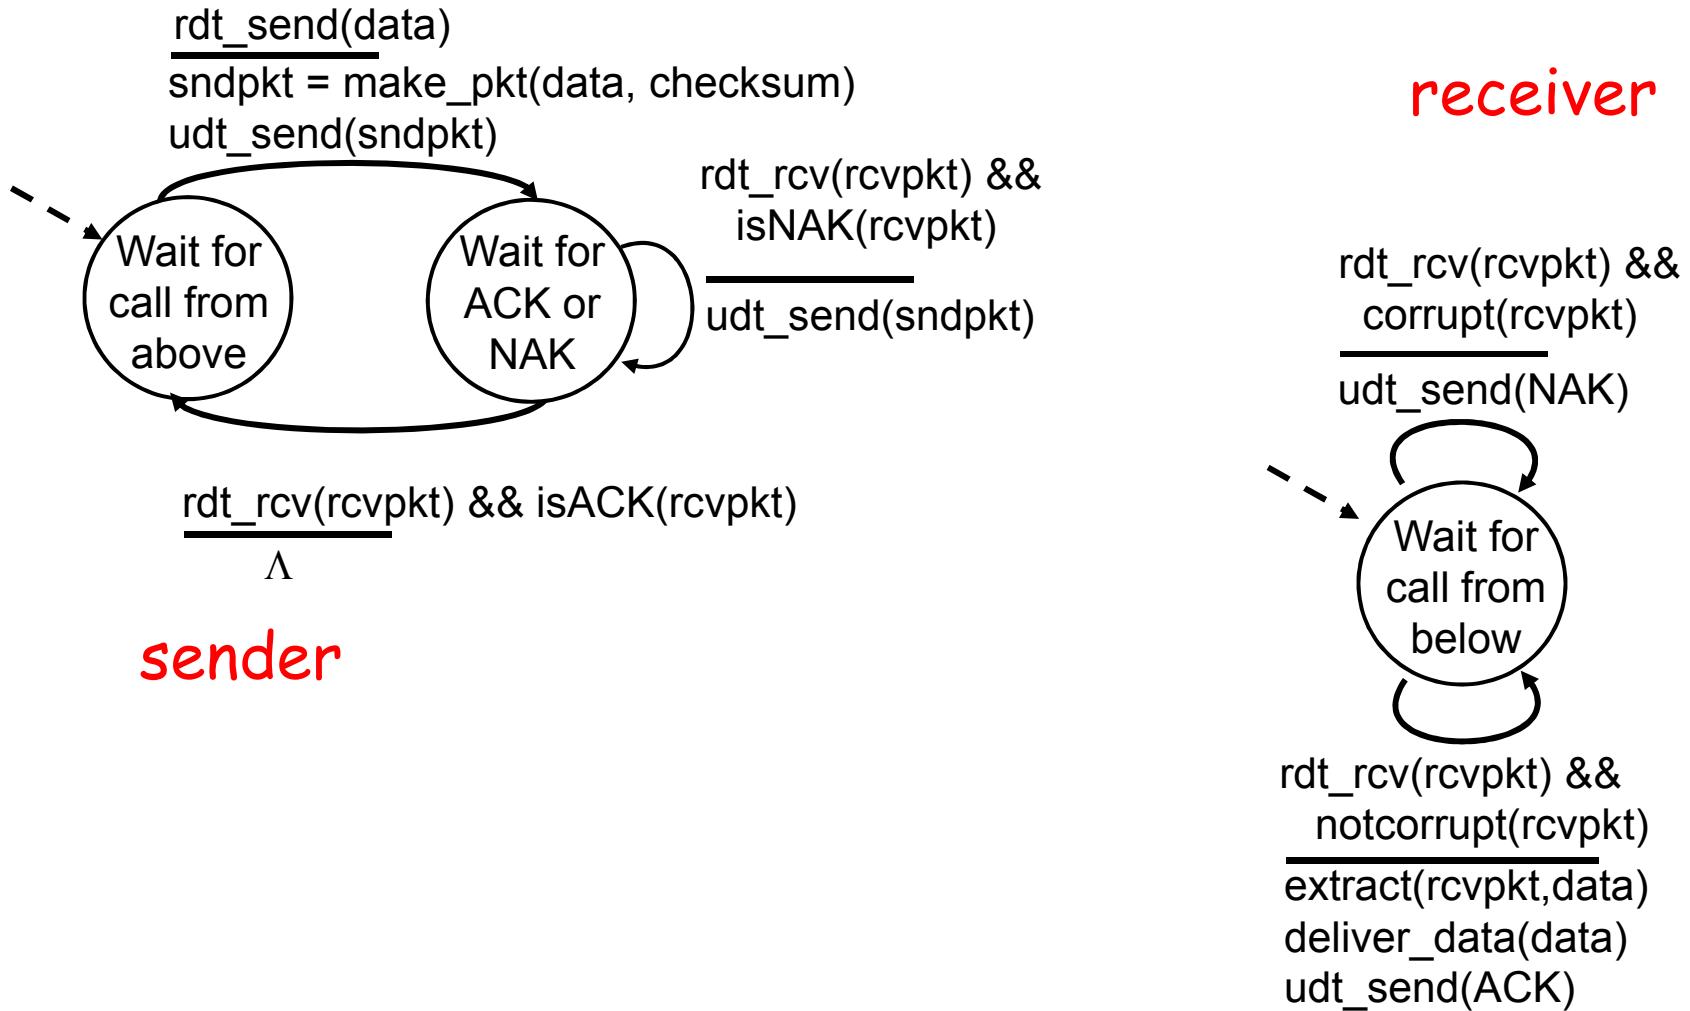
How do humans recover from "errors" during conversation?

List three ways we handle spoken communication in a noisy environment.

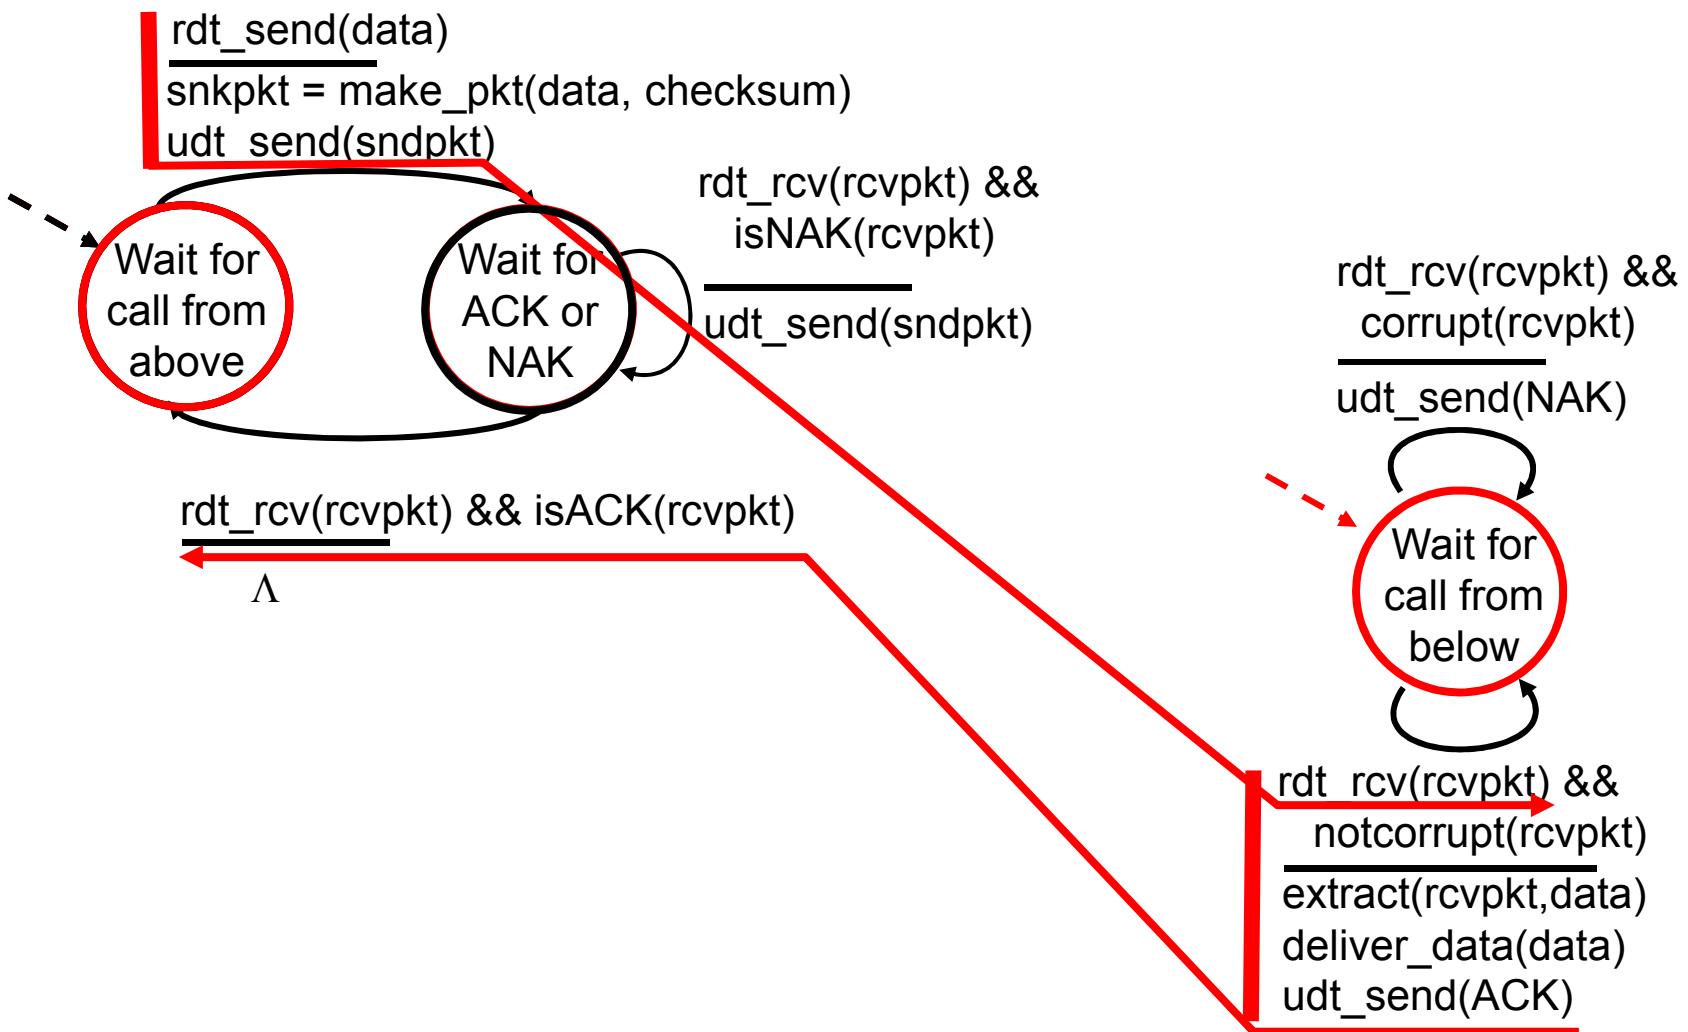
Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ *the question:* how to recover from errors:
 - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- ❖ new mechanisms in `rdt2.0` (beyond `rdt1.0`):
 - error detection
 - receiver feedback: control msgs (ACK,NAK) rcvr->sender
 - retransmission

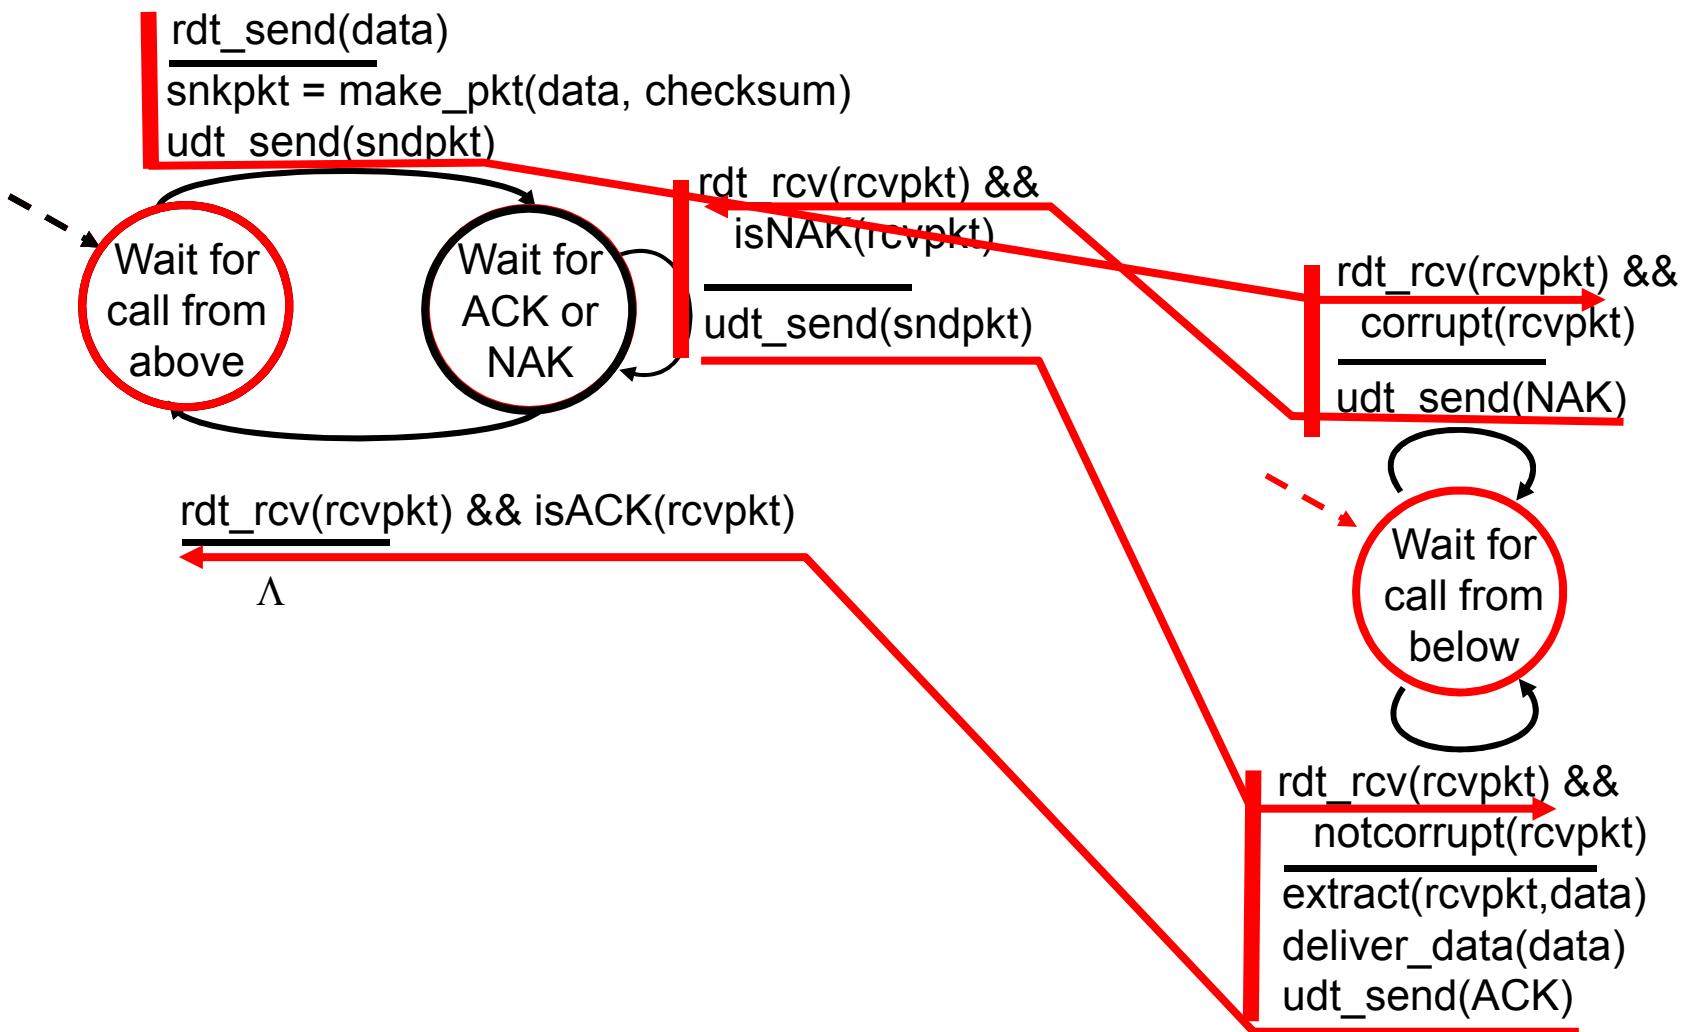
rdt2.0: FSM specification



rdt2.0: operation with no errors



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

- ❖ Can you spot the problem?

rdt2.0 has a fatal flaw!

What happens if
ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit: possible duplicate

❖ How should we handle duplicates?

rdt2.0 has a fatal flaw!

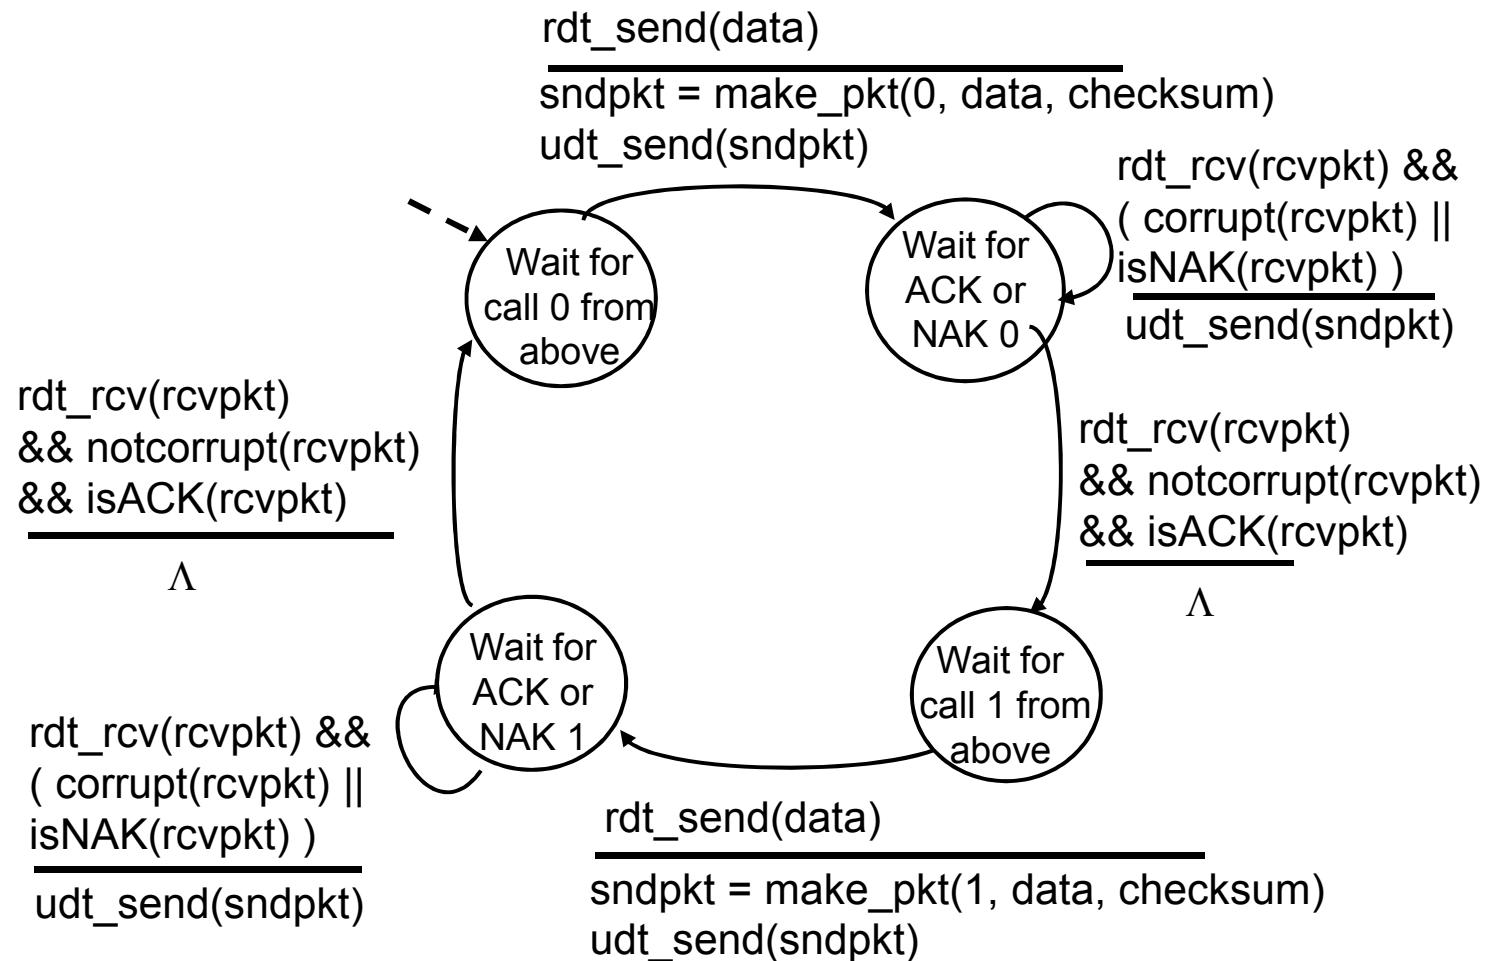
Handling duplicates:

- ❖ sender retransmits current pkt if ACK/NAK garbled
- ❖ sender adds *sequence number* to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkt

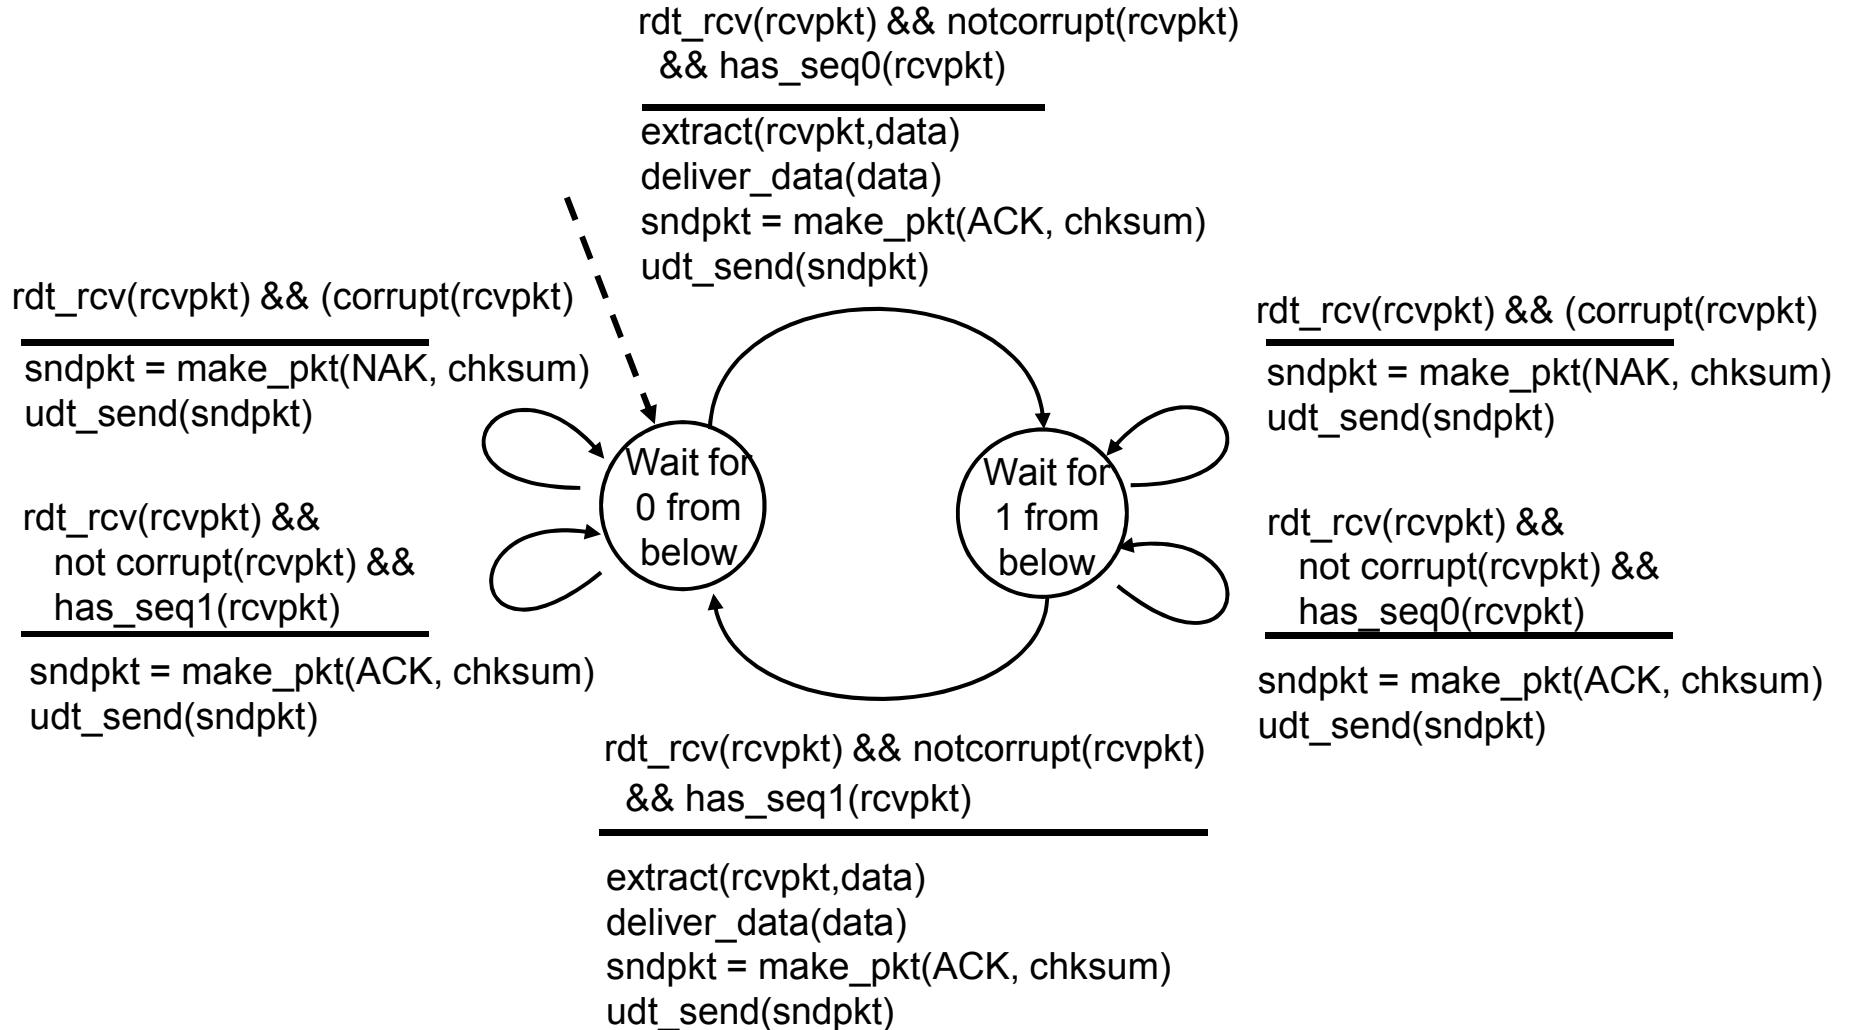
stop and wait

Sender sends one packet,
then waits for receiver
response

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles garbled ACK/NAKs



rdt2.1: discussion

Sender:

- ❖ seq # added to pkt
- ❖ two seq. #'s (0,1) will suffice. Why?
- ❖ must check if received ACK/NAK corrupted
- ❖ twice as many states
 - state must "remember" whether "current" pkt has 0 or 1 seq. #

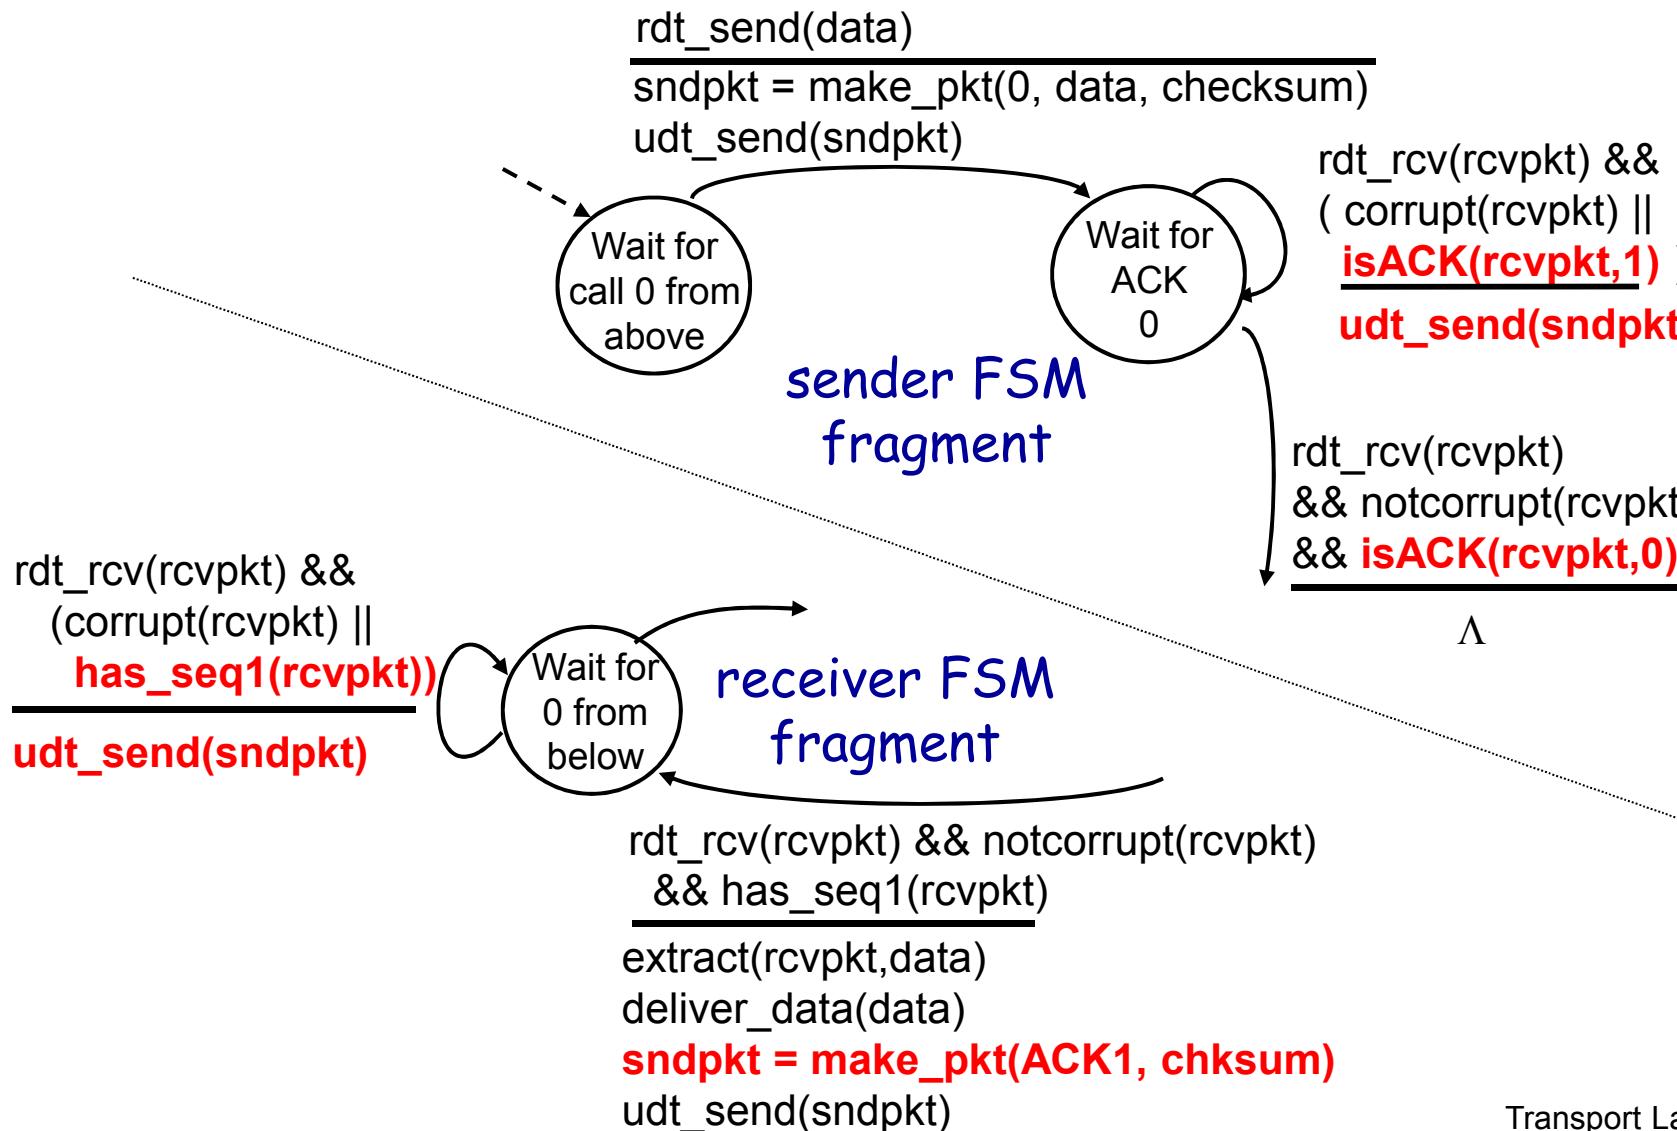
Receiver:

- ❖ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- ❖ note: receiver can *not* know if its last ACK/NAK received OK at sender

rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, using ACKs only
- ❖ instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- ❖ duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: sender, receiver fragments



rdt3.0: channels with errors and loss

New assumption:

underlying channel can also lose packets (data or ACKs)

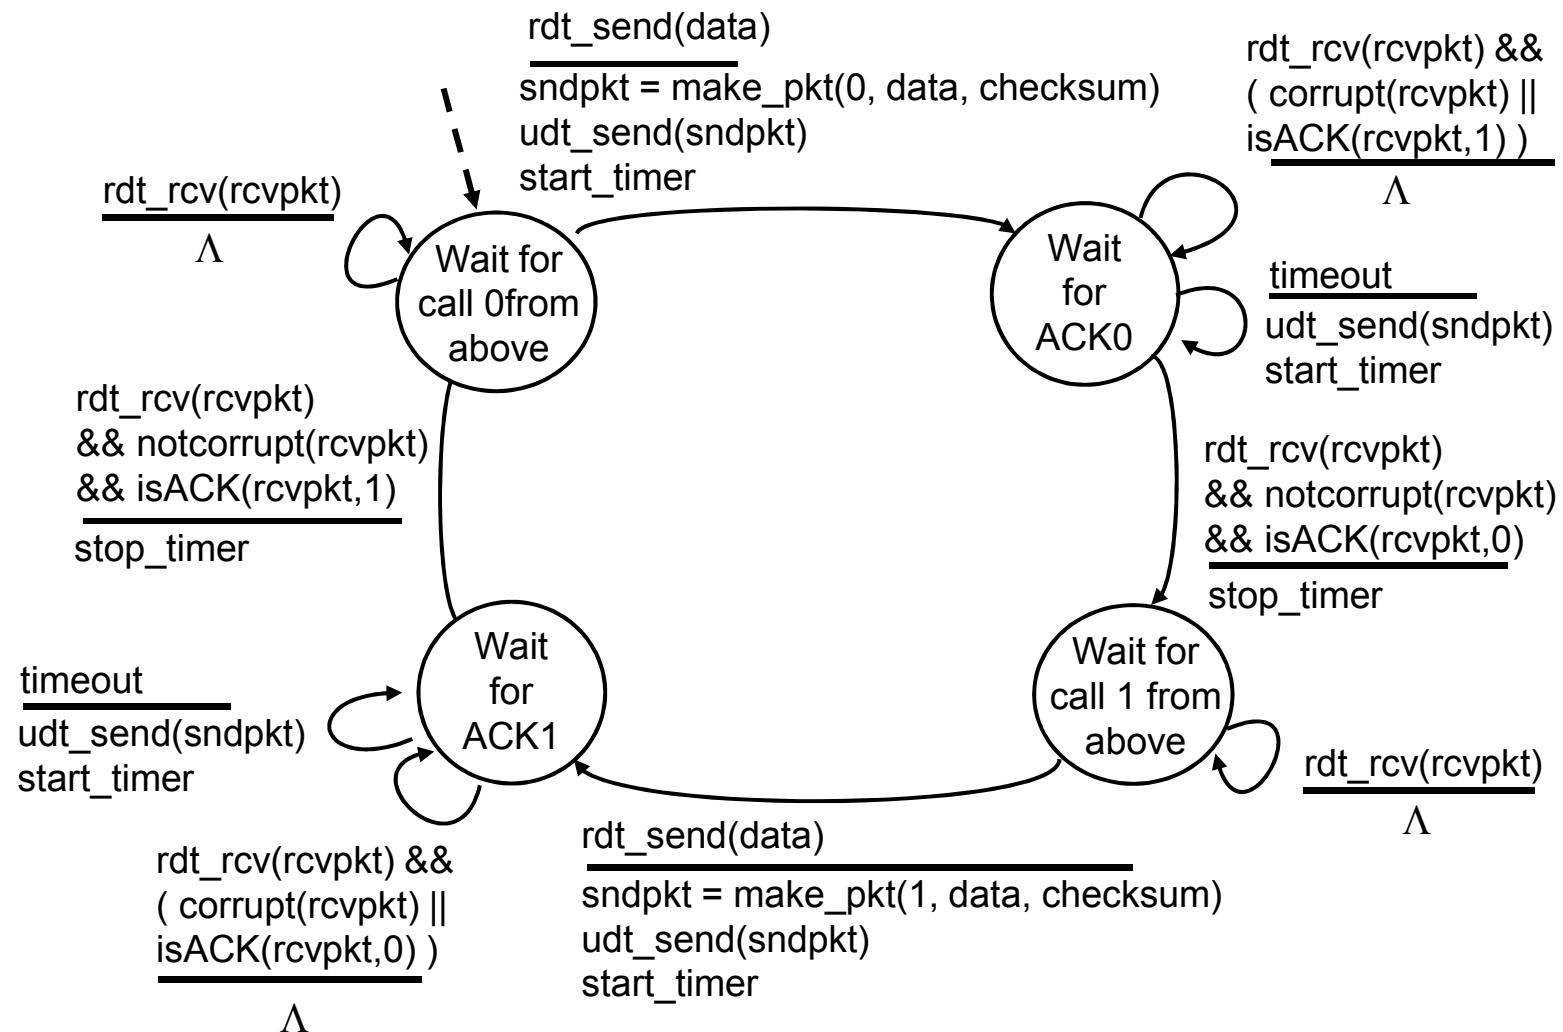
- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

Approach: sender waits

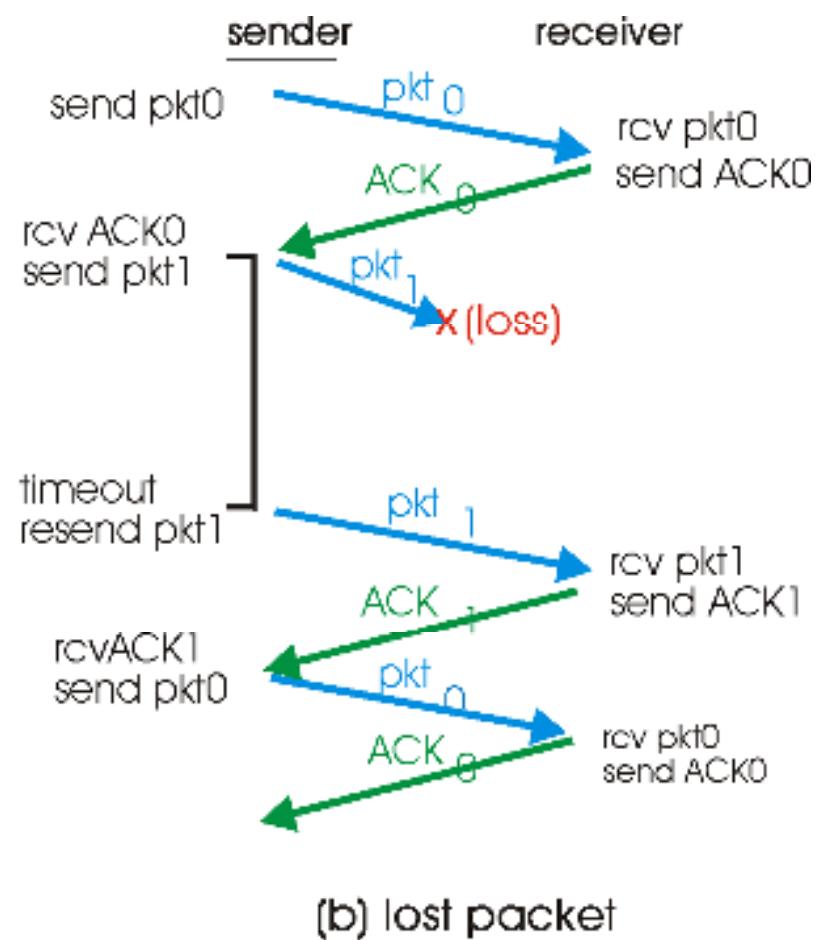
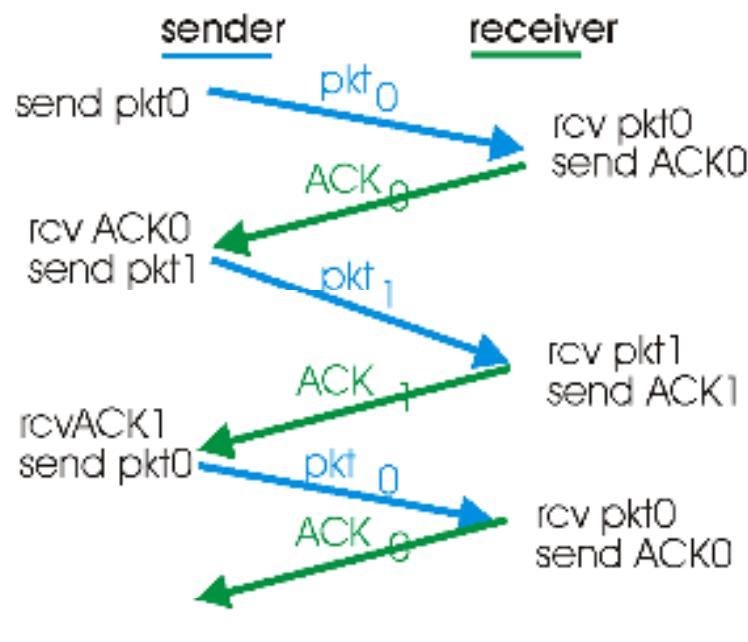
"reasonable" amount of time for ACK

- ❖ retransmits if no ACK received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer

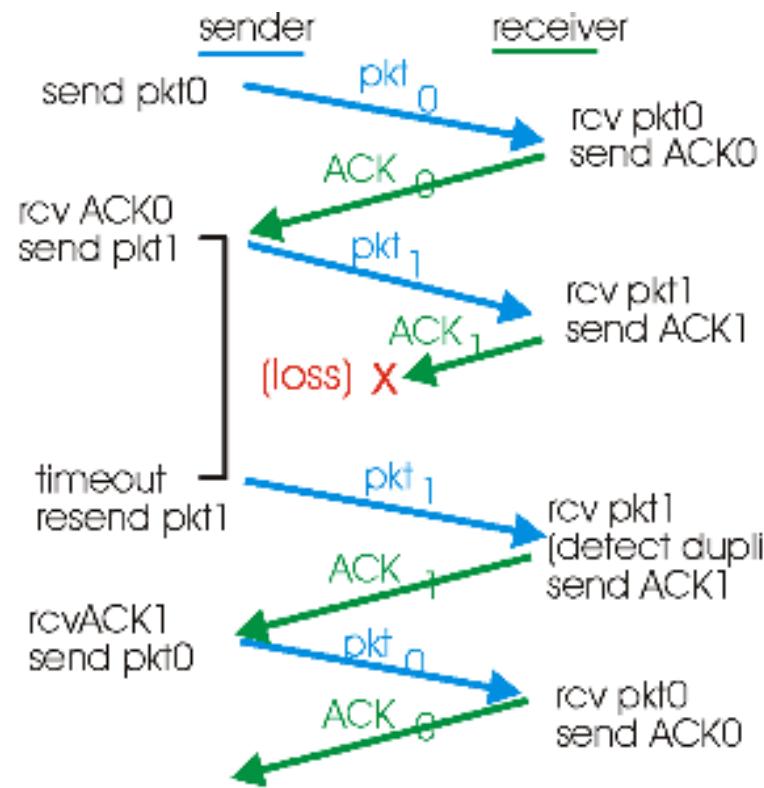
rdt3.0 sender



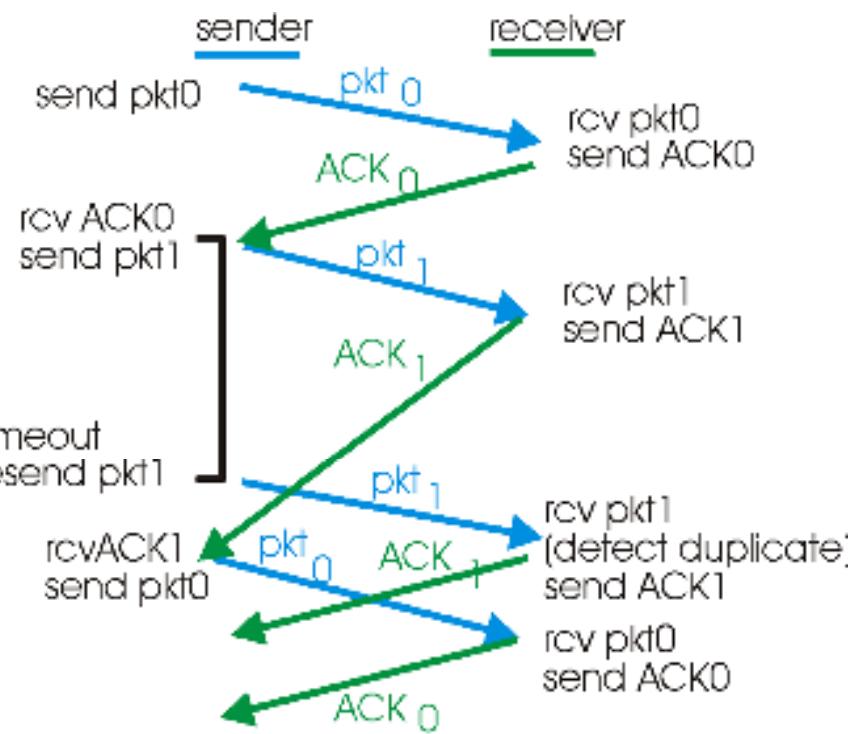
rdt3.0 in action



rdt3.0 in action



(c) lost ACK



(d) premature timeout

Performance of rdt3.0

- ❖ rdt3.0 works, but performance is very poor
- ❖ ex: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

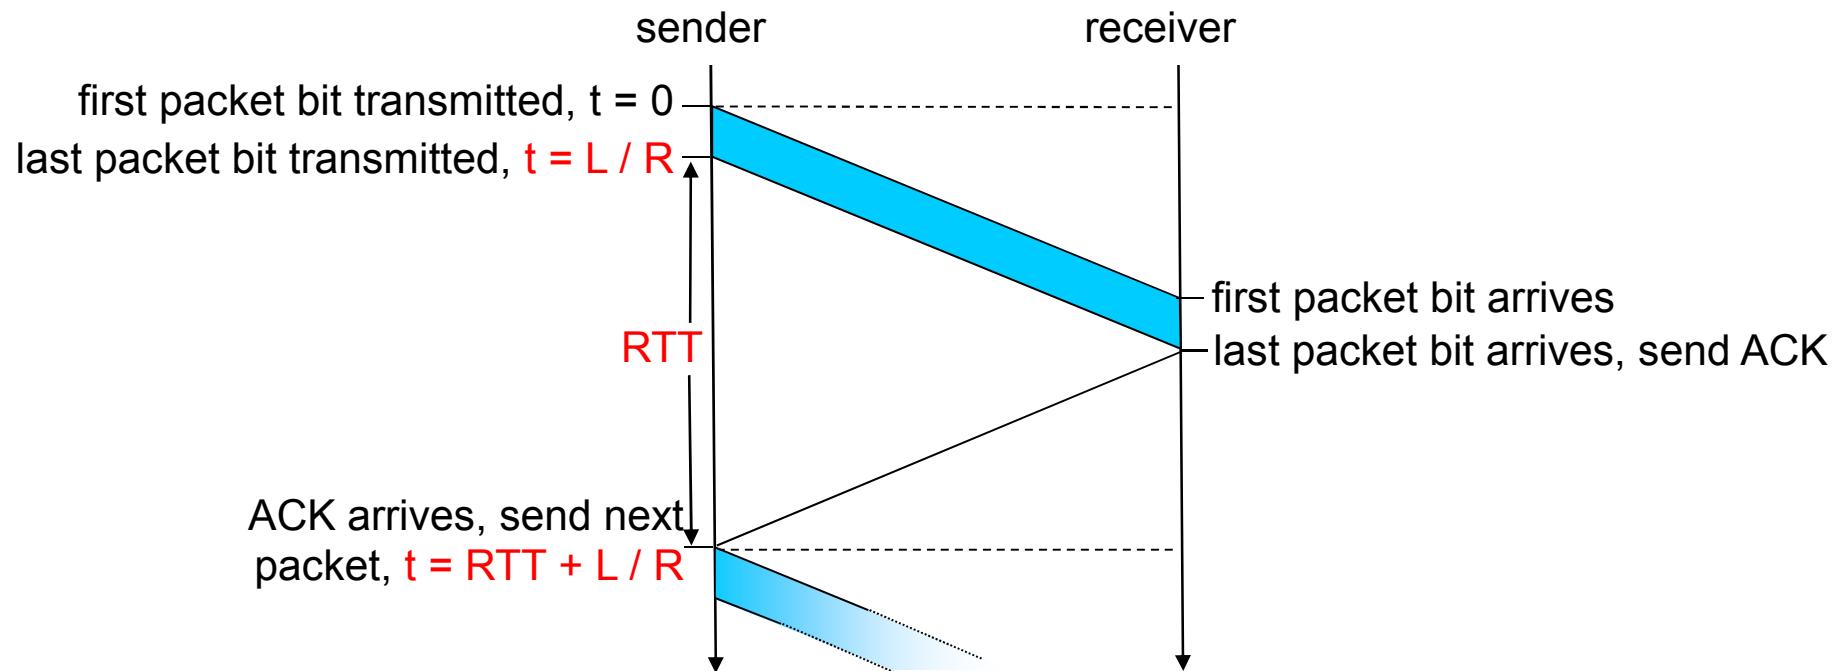
$$d_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bps}} = 8 \text{ microseconds}$$

- U_{sender} : **utilization** - fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

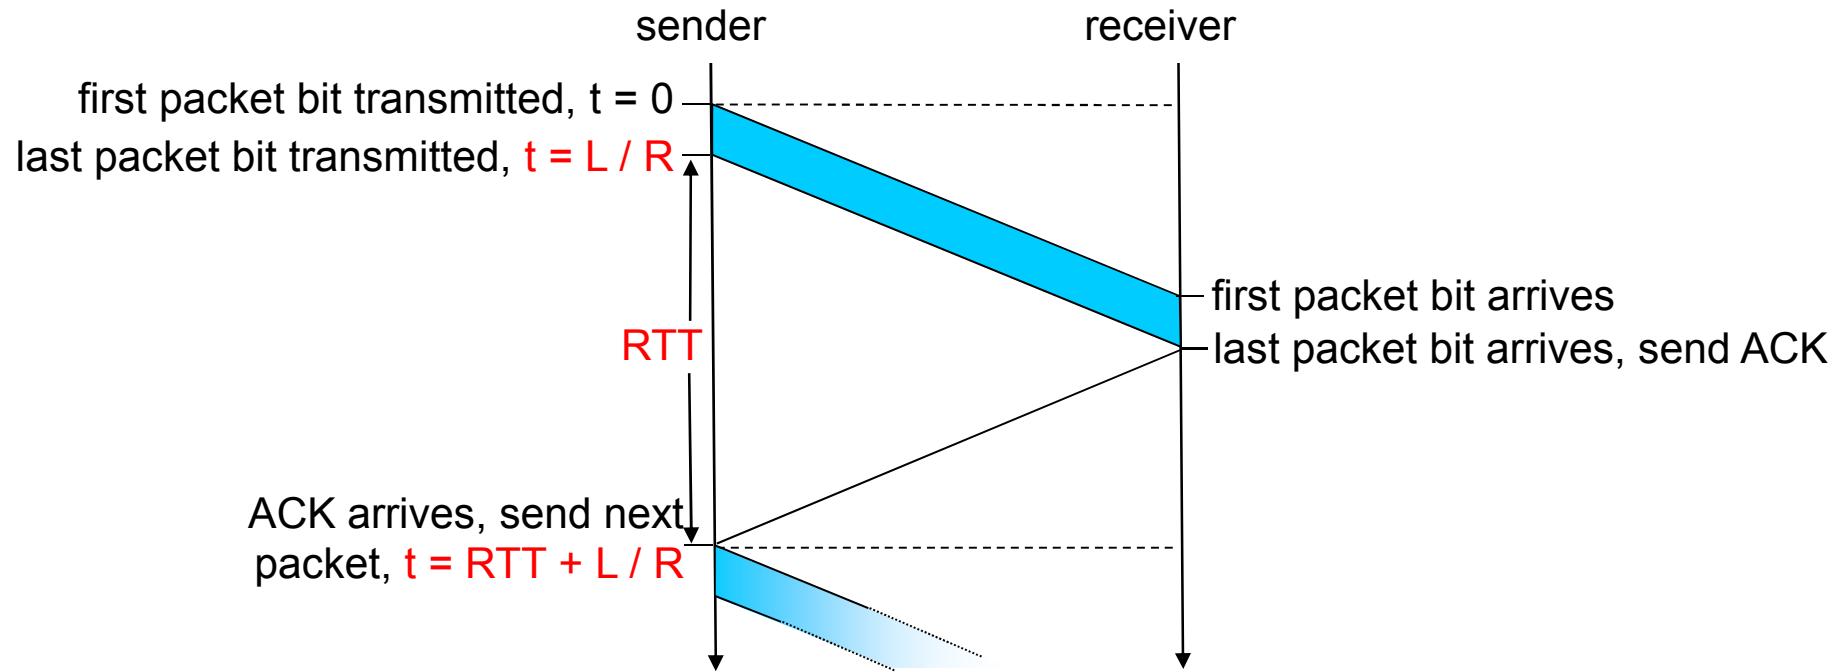
- if RTT=30 msec, 1KB pkt every 30 msec \rightarrow 33kB/sec thruput over 1 Gbps link
- network protocol limits use of physical resources!

rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

rdt3.0: stop-and-wait operation

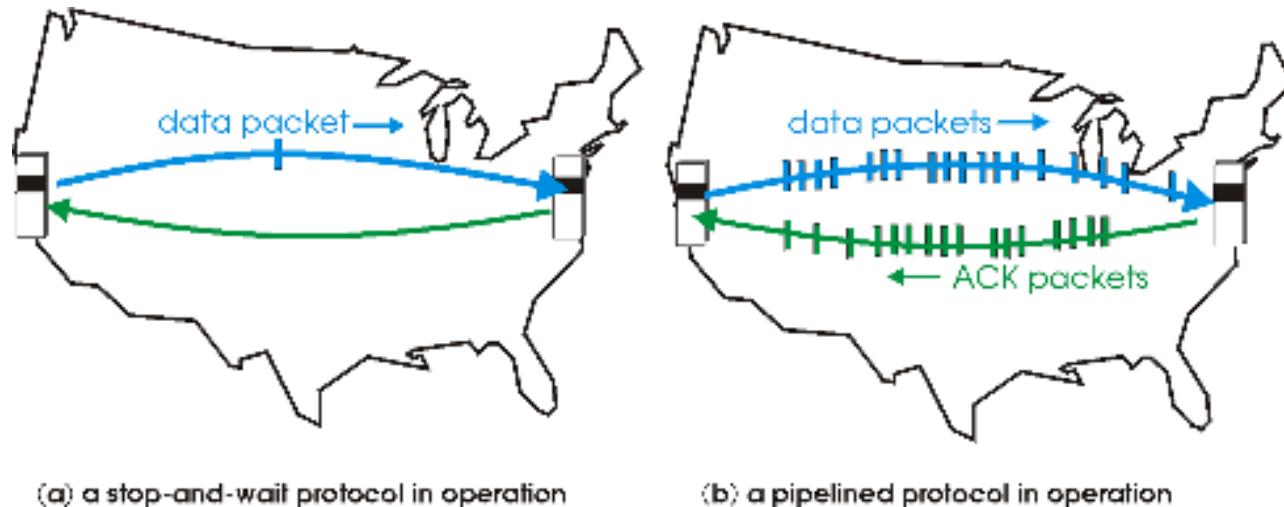


How do we solve the poor performance problem???

Pipelined protocols

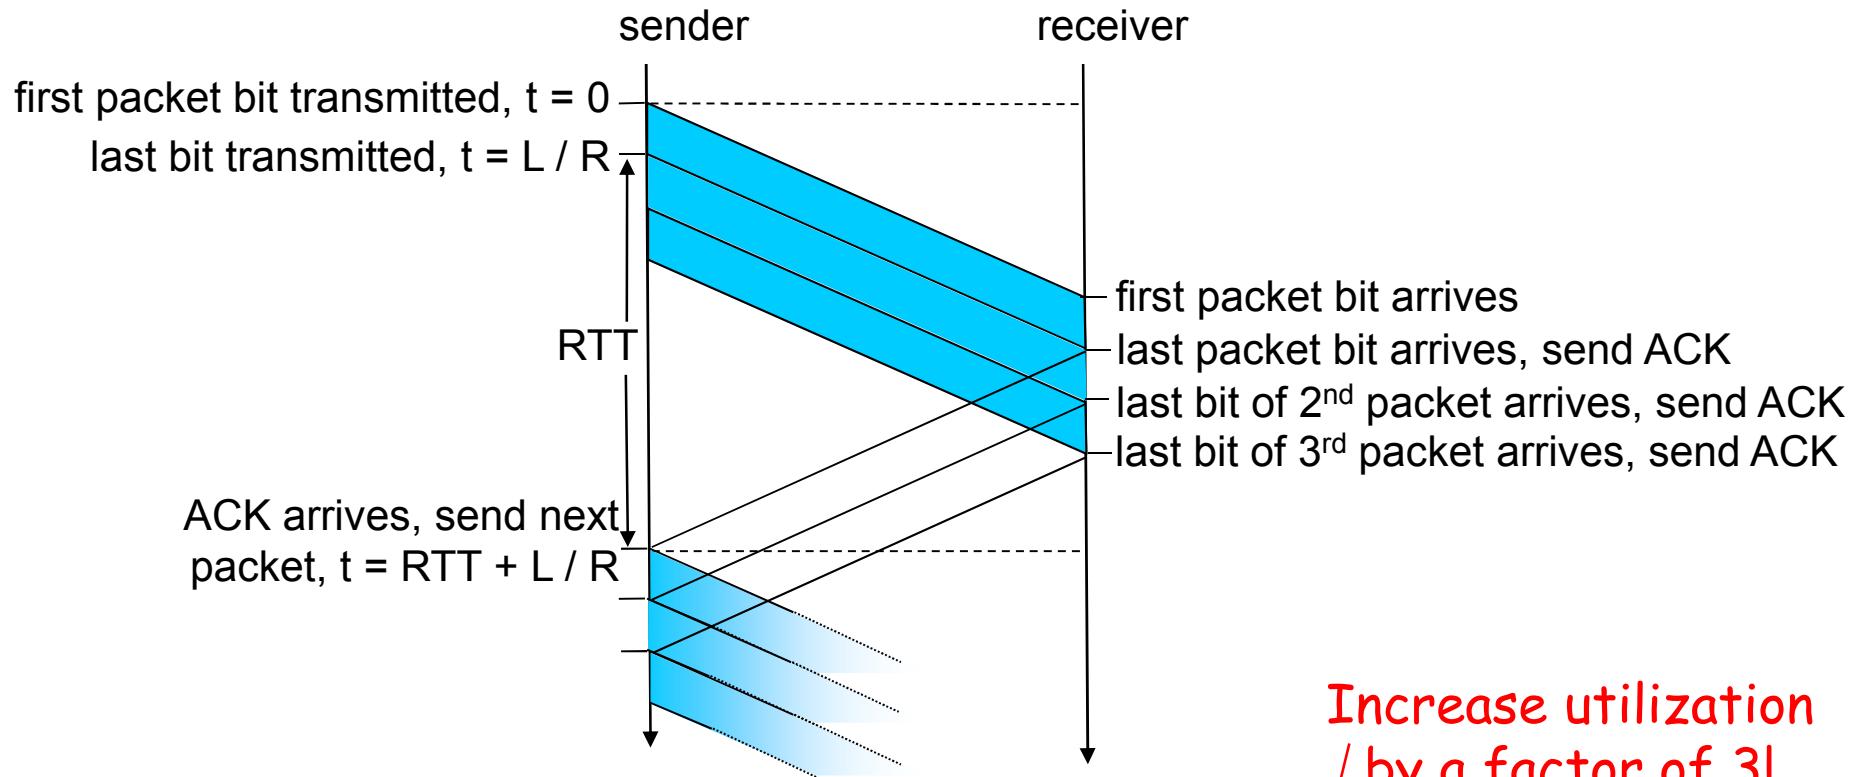
pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver



- ❖ two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Pipelined Protocols

Go-back-N: big picture:

- ❖ sender can have up to N unacked packets in pipeline
- ❖ rcvr only sends *cumulative* acks
 - doesn't ack packet if there's a gap
- ❖ sender has timer for oldest unacked packet
 - if timer expires, retransmit all unack'd packets

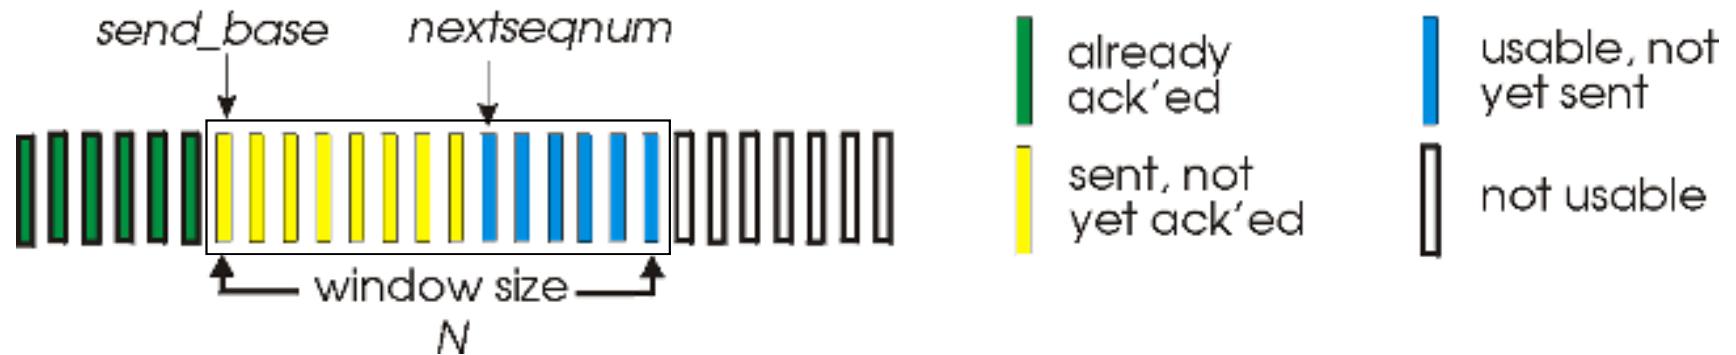
Selective Repeat: big pic

- ❖ sender can have up to N unack'd packets in pipeline
- ❖ rcvr sends *individual ack* for each packet
- ❖ sender maintains timer for each unacked packet
 - when timer expires, retransmit only unack'd packet

Go-Back-N

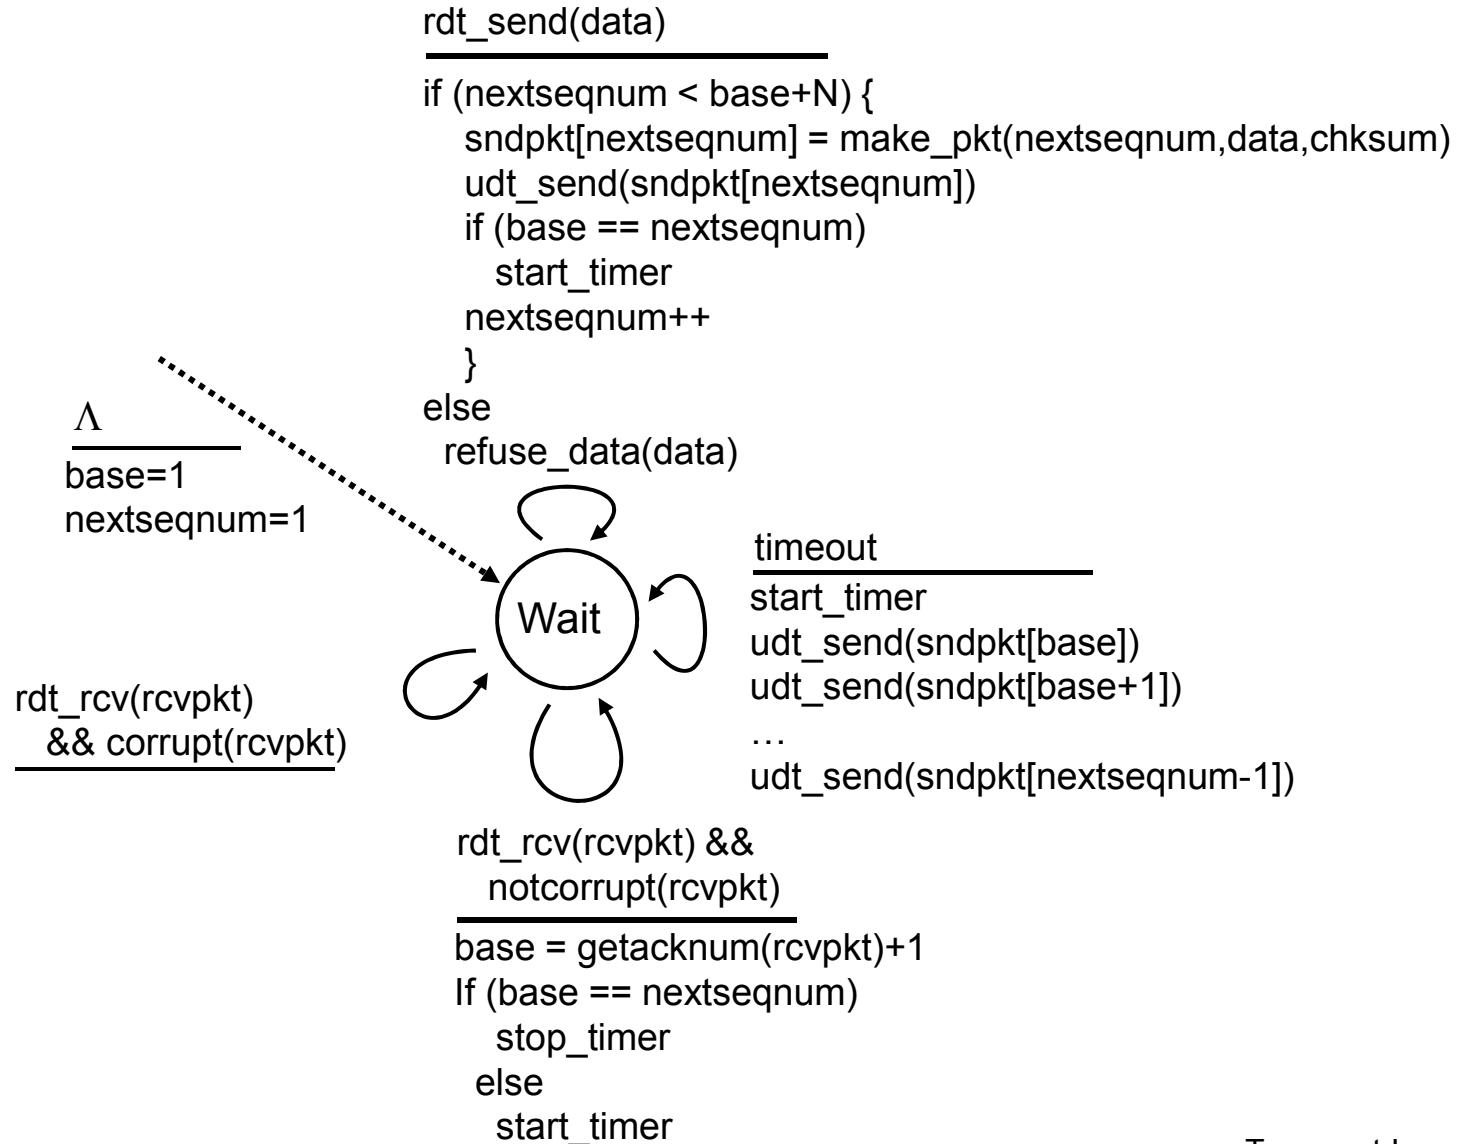
Sender:

- ❖ k-bit seq # in pkt header
- ❖ "window" of up to N, consecutive unack'd pkts allowed

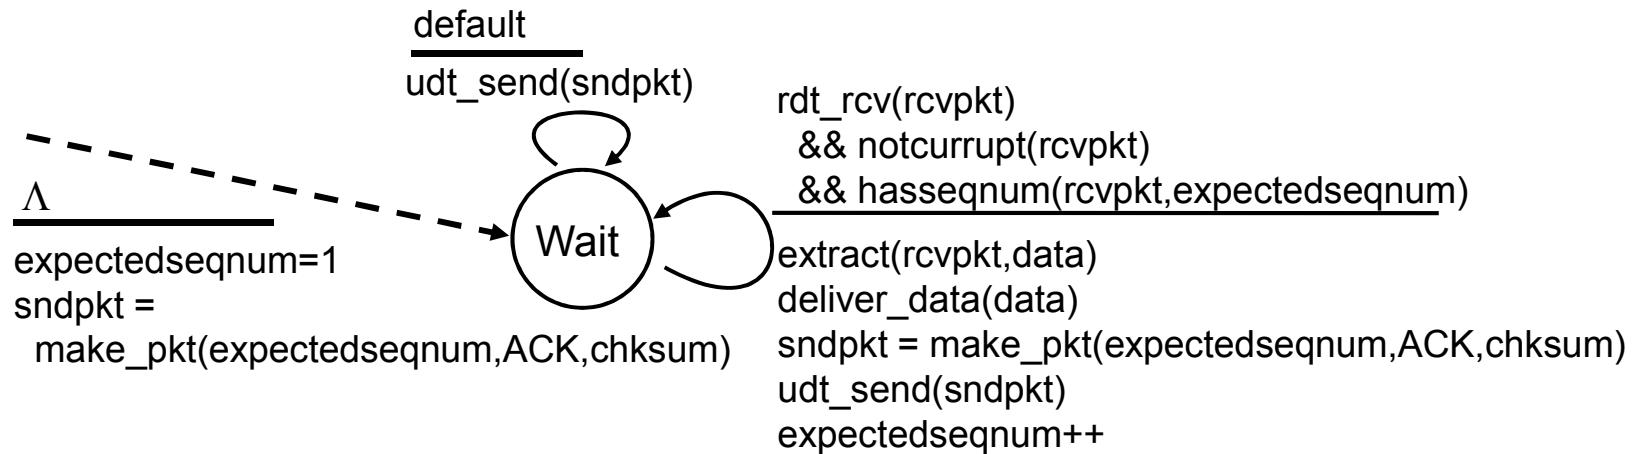


- ❖ ACK(n): ACKs all pkts up to, including seq # n - "cumulative ACK"
 - may receive duplicate ACKs (see receiver)
- ❖ timer for each in-flight pkt
- ❖ *timeout(n)*: retransmit pkt n and all higher seq # pkts in window

GBN: sender extended FSM



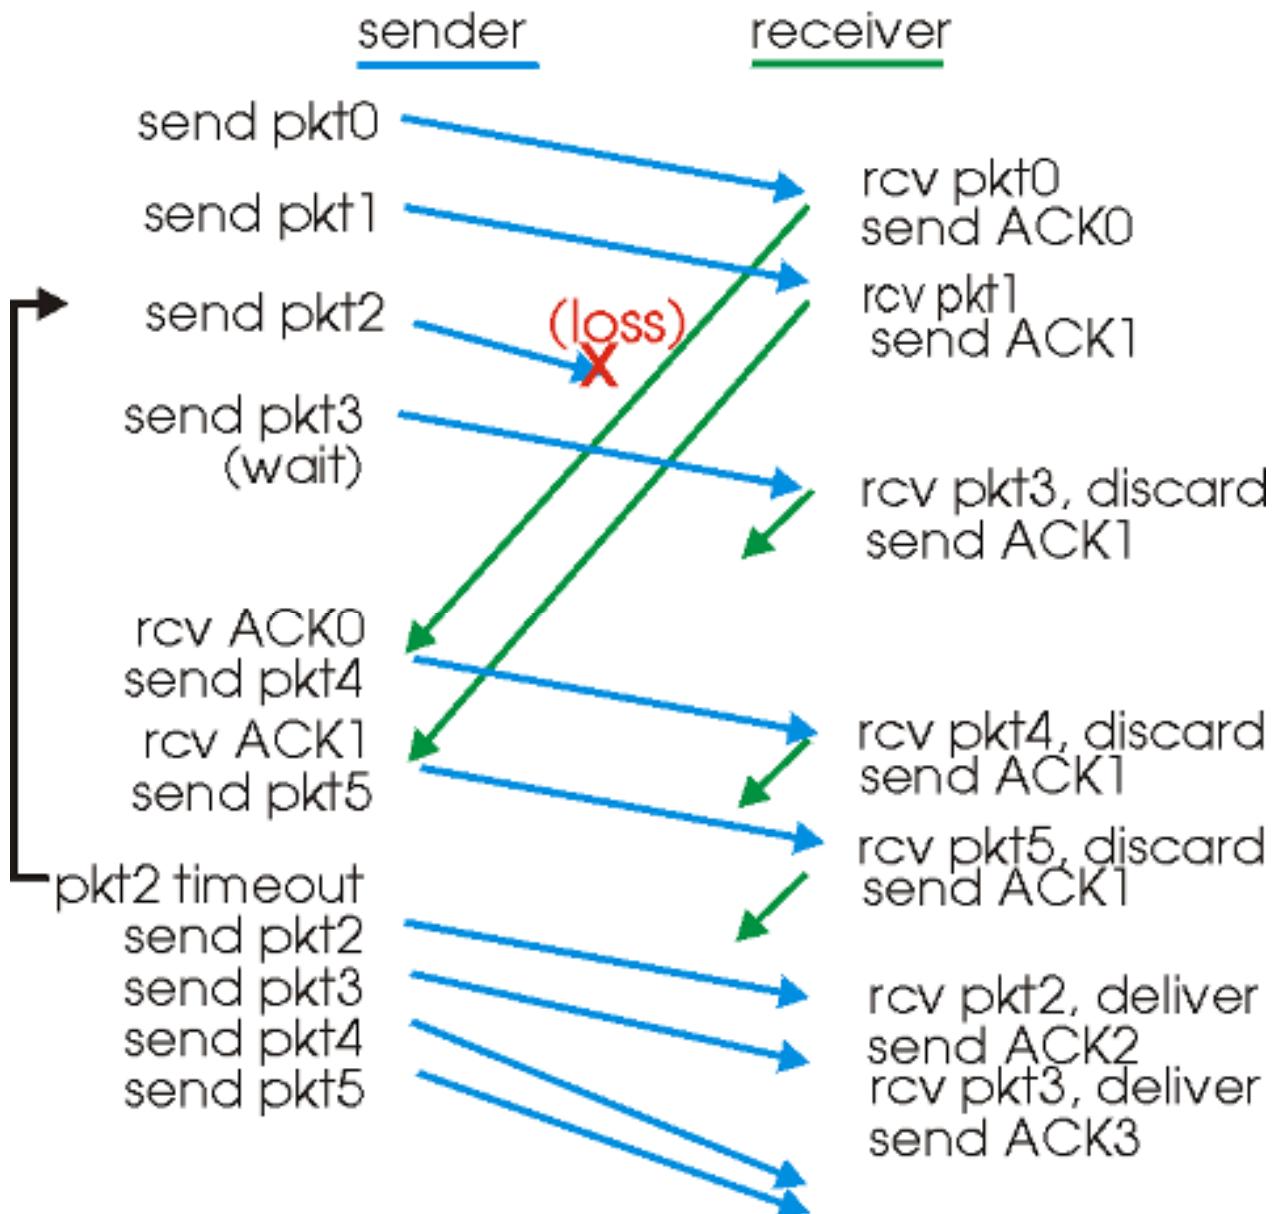
GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- may generate duplicate ACKs
 - need only remember **expectedseqnum**
- ❖ out-of-order pkt:
- discard (don't buffer) -> **no receiver buffering!**
 - Re-ACK pkt with highest in-order seq #

GBN in action



GBN in action:

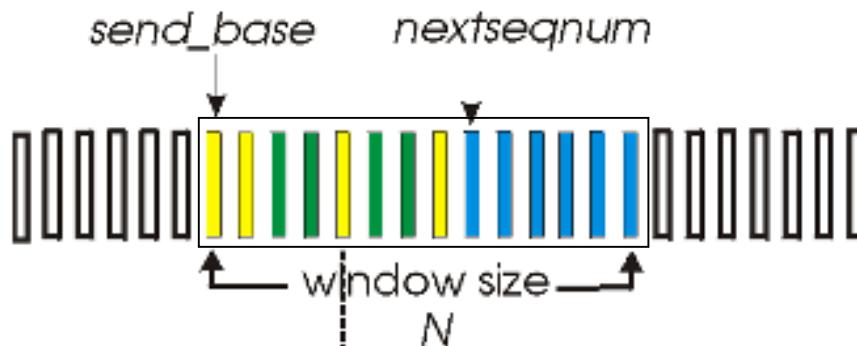
See animated applet at:

http://media.pearsoncmg.com/aw/aw_kurose_network_4/applets/go-back-n/index.html

Selective Repeat

- ❖ receiver *individually* acknowledges all correctly received pkts
 - buffers pkts, as needed, for eventual in-order delivery to upper layer
- ❖ sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt
- ❖ sender window
 - N consecutive seq #'s
 - again limits seq #'s of sent, unACK'ed pkts

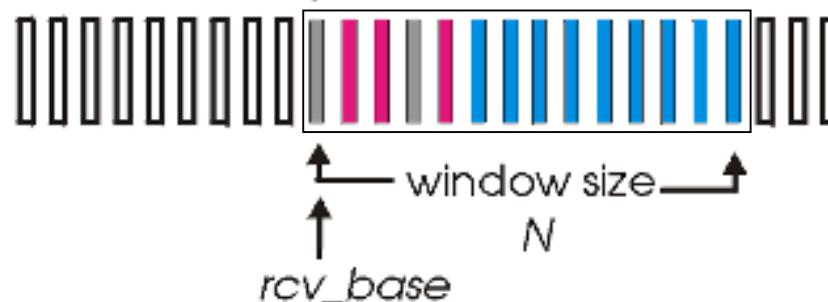
Selective repeat: sender, receiver windows



already
ack'ed
sent, not
yet ack'ed
not usable

usable, not
yet sent
not usable

(a) sender view of sequence numbers



out of order
(buffered) but
already ack'ed
Expected, not
yet received
not usable

acceptable
(within window)
not usable

(b) receiver view of sequence numbers

Selective repeat

sender

data from above :

- ❖ if next available seq # in window, send pkt

timeout(n):

- ❖ resend pkt n, restart timer

ACK(n) in [sendbase,sendbase+N]:

- ❖ mark pkt n as received
- ❖ if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- ❖ send ACK(n)
- ❖ out-of-order: buffer
- ❖ in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

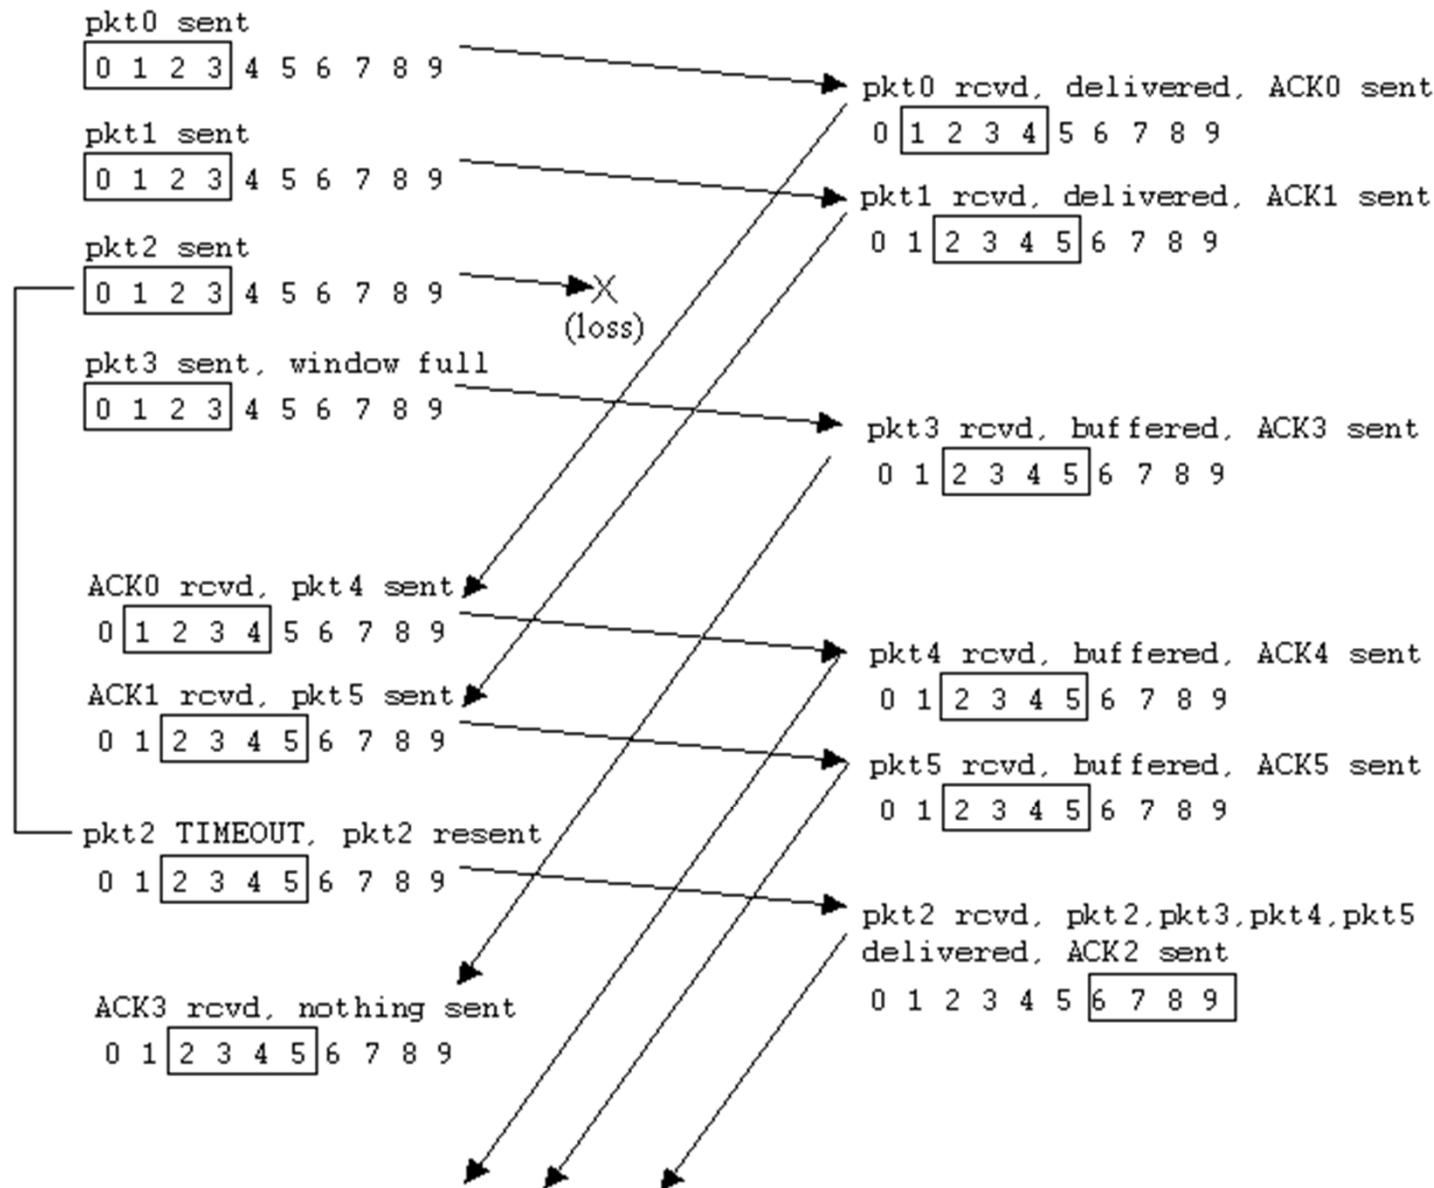
pkt n in [rcvbase-N,rcvbase-1]

- ❖ ACK(n)

otherwise:

- ❖ ignore

Selective repeat in action

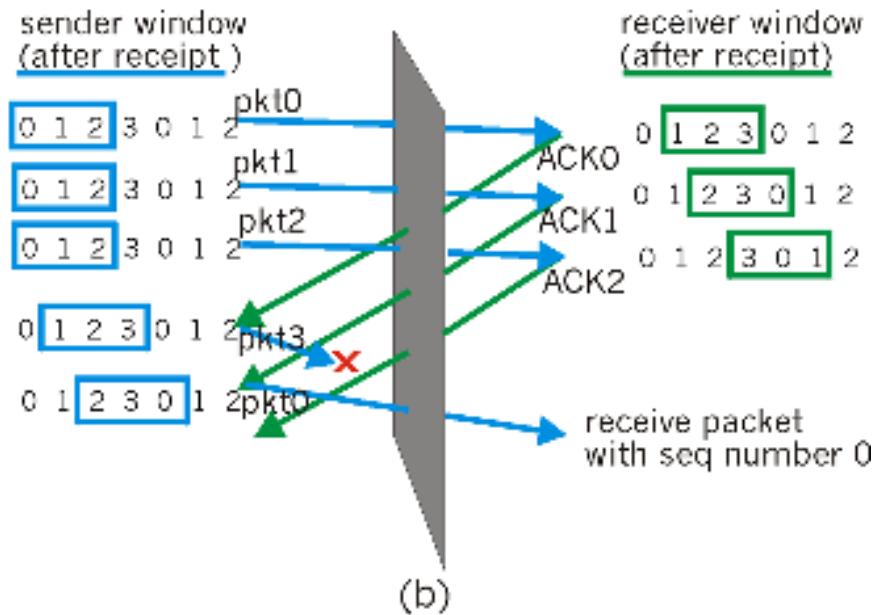
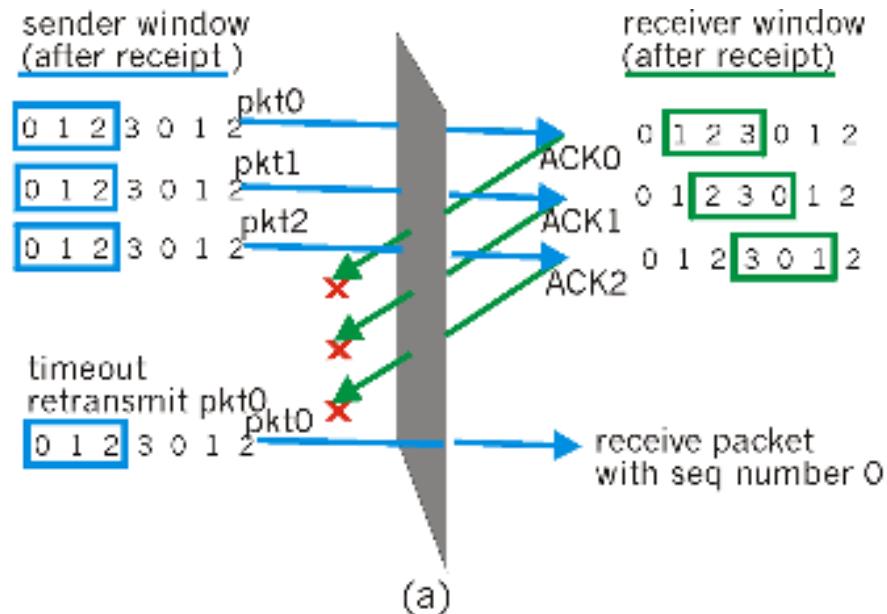


Selective repeat: dilemma

Example:

- ❖ seq #'s: 0, 1, 2, 3
- ❖ window size=3
- ❖ receiver sees no difference in two scenarios!
- ❖ incorrectly passes duplicate data as new in (a)

Q: what relationship between seq # size and window size?



Selective repeat:

- ❖ See animated applet:
- ❖ http://media.pearsoncmg.com/aw/aw_kurose_network_4/applets/SR/index.html

Lesson 9: Summary

- ❖ Reliable data transfer is an optional transport layer service.
- ❖ Reliable data transfer using the stop-and-wait-for-ack approach has very poor performance.
- ❖ Pipelining addresses stop-and-wait performance issues. Two generic pipelining approaches:
 - Go-Back-N
 - Selective Repeat.
- ❖ We will see in the next lesson how TCP solves the RDT problem with a pipelined approach that is a combination of Go-Back-N and Selective Repeat.